

RevIN Experiment 3: Centering and Transform Variants

1 Theoretical overview and goal

Mean centering assumes the look-back average is the appropriate location anchor for the future. Last-value centering instead predicts changes around the most recent observation, which may better suit local trends. The inverse hyperbolic sine transform

$$z = \operatorname{asinh}(x)$$

is approximately linear near zero and logarithmic for large magnitudes, giving a reversible compression of heavy-tailed values before instance normalization.

Experimental goal. Test two lightweight changes to RevIN’s inductive bias: a more local location anchor and a nonlinear magnitude compression as appendix ablations.

2 Implementation details

The variants use non-affine instance normalization and normalized MSE:

Method	Center	Transform
<code>instance_nmse</code>	Look-back mean	Identity
<code>revin_last_nmse</code>	Last observation	Identity
<code>revin_arcsinh_nmse</code>	Look-back mean	<code>asinh</code>

Every transformation is inverted before data-space evaluation. These appendix variants run in the full and ultra profiles, over the same six datasets, six settings, PatchTST (plus DLinear in ultra), three seeds, optimizer, batch size, 10,000-step budget, and validation frequency. Configuration mapping is defined by `method_args` in `src/slurm/run_revin_experiment.sh`; layer behavior is in `src/utlis/models.py`.

Tables are separate by model and test split. They report ordinary MSE as mean $\pm$ sample standard deviation with explicit $\times 10^m$ row scaling. The baseline for both contrasts is `instance_nmse`.

3 Interpretation

Last centering tests sensitivity to local level evolution; `asinh` tests sensitivity to tails and multiplicative scale. A gain is evidence for that specific preprocessing bias, not a general solution to conditional shift. Results should be examined across datasets because either transform may remove useful absolute-level information.